

# Fitting the Undetected: Light-curve Feature Estimation in the Presence of Non-detections

Some wild & crazy guys

## ABSTRACT

In multi-epoch imaging surveys, faint variable sources will not be detectable at all epochs. In principle, the best approach is to perform photometric measurements anyway, so that there is a datum required at every epoch. In practice, in most surveys we are provided at some epochs only with upper limits; in some not even upper limits but just the information that the star was not observed. Here we demonstrate with real data on periodic variables from the All-Sky Automated Survey (ASAS) that these non-detections are nonetheless useful in fitting and parameter estimation; their inclusion improves both precision and accuracy. One novel aspect of this work is that we do not require the data source to include accurate—or even any—information about the uncertainty variances on detections or the values of the upper limits; we find that we can model these as latent variables, even when they are different at every epoch. Using realistic simulations of Mira variable stars, we obtain amplitude estimates that are **XXX%** more accurate by incorporating the non-detections into the analysis. For Mira and RR Lyrae variables observed by ASAS, our method obtains a tighter period - amplitude relationship than the standard technique.

## 1. Introduction

Surveys working in real-time tend to analyze each image with an eye to populating databases quickly. Unless a particular region of the sky gets special treatment (such as it is a known place of interest), discovery of variability happens by comparison of flux values at the database level. Analysis on such databases of transient/variables thus deals with censored data where there is often less information available than there should be. Even then, we might be able to infer what the thresholds were based on S/N measurements of objects neighboring the source of interest.

♣ **joey:** What people actually do, what problems that could cause, and what people should do ♣

A photometric survey (think of it as operating in a single photometric bandpass for now) scans over the celestial position of a particular variable star a large number of times  $t_i$ . At each of these times, the imaging data are analyzed by software, which treats each scan as a *completely independent* survey. That is, when analyzing the data from time  $t_i$ , none of the information from any of the other times is used in any way.

At each time  $t_i$ , the star is either detected ( $q_i = 1$ ) by the software or not ( $q_i = 0$ ). When it is detected, the software returns a (possibly bad) flux value  $f_i$  and (likely bad) uncertainty variance  $s_i^2$ . When the source is *not* detected ( $q_i = 0$ ), nothing is reported—in this case, because the software is not awesome, it doesn’t see any reason to report anything at all, since on the non-detect passes, it has no inkling that this patch of the sky is interesting in any way. Welcome to the desert of the real.

Beyond this, there are two additional issues. The first is that each night of imaging is different in an unreported and unknown way. That is, there is different transparency, sky brightness, and point-spread function. So the detection limit, or completeness level, or censoring of the catalog is different every night. The second issue is that either because we are getting the data from a non-generous source, or because the conditions change rapidly and unpredictably, we can’t analyze all the sources in a finite patch simultaneously. For each variable star of interest, we *only* get the data on that star itself.

♣ **joey:** RRL end at 120 kpc in Sesar et al. paper ♣

♣ **joey:** discuss in context of: parameter estimation (amplitude / period), discovery (we’re gonna find a lot more variable sources if we take the non-detections into account), and classification; lossiness ♣

♣ **joey:** Numbers of LCs affected: Stripe 82, ASAS ♣

## 2. Method

To recap, all we get at each epoch  $t_i$  is a bit  $q_i$  and, when  $q_i = 1$ , two data  $f_i$  and  $s_i^2$ . We don’t get any information about the upper limits or detection thresholds at the times  $t_i$  when  $q_i = 0$ , and we don’t know anything about the data quality (catalog veracity) at any epoch. For notational convenience we assemble all the data into an enormous set  $D$  given by

$$D \equiv \{D_i\} \tag{1}$$

$$D_i \equiv (q_i, f_i, s_i^2) \quad , \quad (2)$$

where the  $f_i$  and  $s_i^2$  data will be randomly generated or meaningless at times at which the  $q_i = 0$ .

We don't have access to all the information we need to use responsibly the detections and non-detections in model fitting and inference. We have to instantiate latent variables to stand in for the missing information. For example, since we don't know the censoring threshold or completeness limit for any observation  $i$ , we have to make it a model parameter  $b_i$ . Because we don't believe the reported uncertainty variances  $s_i^2$ , we have to add parameters that model the departure of the true uncertainty variance  $\sigma_i^2$  from the reported variance, and so on.

For the limited purposes here, the variable star will be periodic with angular frequency  $\omega$ . The expected brightness  $\mu_i$  of the star at time  $t_i$  will be given by a linear combination of periodic functions something like

$$\mu_i = \sum_k A_k x_k(t_i|\omega) \quad , \quad (3)$$

where the  $A_k$  are coefficients and the  $x_k$  are functions. This can be generalized easily. In addition to this expected brightness, we assume that there is some either model uncertainty or stochastic variation, leading to a *model variance*  $s_\mu^2$  at each point (assumed constant but easily generalized).

If a source is detected at epoch  $t_i$ , we assume that its observed flux  $f_i$  is related to the expected flux  $\mu_i$  by a Gaussian with variance that has measurement and model contributions

$$p(f_i|q_i, \sigma_i, \theta, I) = \begin{cases} 1 & \text{for } q_i = 0 \\ N(f_i|\mu_i, s_\mu^2 + \sigma_i^2) & \text{for } q_i = 1 \end{cases} \quad (4)$$

$$\theta \equiv (\omega, \{A_k\}, s_\mu^2, \dots) \quad (5)$$

$$I \equiv (\{t_i\}, \text{assumptions}) \quad , \quad (6)$$

where we have made the prior information about observation times  $t_i$  and other assumptions explicit in the prior information blob  $I$ , and we have a parameter blob  $\theta$ , which we will specify completely below. Note that the flux datum is ignored when the source is undetected.

Similarly for the reported uncertainty variance  $s_i^2$ . We don't assume that it is correct in any case, but when the source is detected, we assume that it is connected to the true uncertainty  $\sigma_i$  by a probabilistic likelihood, also Gaussian

$$p(s_i|q_i, \sigma_i, \theta, I) = \begin{cases} 1 & \text{for } q_i = 0 \\ N(s_i|\sigma_i, v_\sigma) & \text{for } q_i = 1 \end{cases} \quad (7)$$

$$\theta \equiv (\omega, \{A_k\}, s_\mu^2, v_\sigma, \dots) \quad , \quad (8)$$

where we have added a model parameter  $v_\sigma$ , which represents the variance in the distribution of reported uncertainties given the true uncertainty; implicitly  $v_\sigma$  is a component of  $\theta$ . Again, we are ignoring the variance datum when the source is undetected.

Under the assumption that the inclusion (or not) of the source in the catalog is based on (or very strongly related to) the measured flux (rather than the true flux), the natural likelihood for the bit  $q_i$ , on which the above likelihoods are conditioned is

$$P(q_i|b_i, \sigma_i, \theta, I) = \begin{cases} \int_{-\infty}^{b_i} N(f|\mu_i, s_\mu^2 + \sigma_i^2) df & \text{for } q_i = 0 \\ \int_{b_i}^{\infty} N(f|\mu_i, s_\mu^2 + \sigma_i^2) df & \text{for } q_i = 1 \end{cases}, \quad (9)$$

where the integrals end at an observed-flux limit  $b_i$ , which (as mentioned above) is also a model parameter. The motivation for this expression is that it imagines all possible observations  $f$  of the star at time  $t_i$  given the model and uncertainties, and finds the fraction that would make the (imagined) detection cut.

We are going to product together these likelihoods (for  $q_i$ ,  $f_i$ , and  $s_i^2$ ); this looks like it is using the flux twice, once in the PDF for  $f_i$  in expression (4) and once in the probability for  $q_i$  in expression (9). It is not, however, because in the probability for  $q_i$  the flux  $f$  that appears is an integration variable, not an observation or datum. We treat the detections or non-detections as being “prior” to the measurement. This is the weakest plank in our otherwise impenetrable fortress of solitude.

We are going to treat the true detection limit  $b_i$  for each observation, and its true uncertainty  $\sigma_i$ —two parameters *per data point*—as nuisance parameters. We want to marginalize them out. This requires priors on these. Again we choose Gaussians for simplicity, but let the means and variances of these Gaussians be new model parameters—hyperparameters if you wish.

$$p(D_i|\theta, I) = \int P(D_i|b_i, \sigma_i, \theta, I) p(b_i|\theta) p(\sigma_i|\theta) db_i d\sigma_i \quad (10)$$

$$p(D_i|b_i, \sigma_i, \theta, I) = P(q_i|b_i, \sigma_i, \theta, I) p(f_i|q_i, \sigma_i, \theta, I) p(s_i|q_i, \sigma_i, \theta, I) \quad (11)$$

$$p(b_i|\theta) = N(b_i|B, V_B) \quad (12)$$

$$p(\sigma_i|\theta) = N(\sigma_i|S, V_S) \quad (13)$$

$$\theta \equiv (\omega, \{A_k\}, s_\mu^2, v_\sigma, B, V_B, S, V_S) \quad , \quad (14)$$

where we have (at last) explicitly assembled all the variable-star parameters and hyperparameters into the big parameter vector  $\theta$ . The huge sets of nuisance parameters  $\{b_i\}$  and  $\{\sigma_i\}$  don’t remain, because by integration we removed them.

Finally and it goes without saying that if we treat the data points as independent (as we are free to do, given the non-awesomeness of the software data source), we have that

$$p(D|\theta, I) = \prod_i p(D_i|\theta, I) \quad . \quad (15)$$

This is the likelihood for variable-star parameters given the entire data set (all the measurements and non-detections of this star from all the epochs, as delivered by the untrustworthy robots). It “correctly” or at least “justifiably” uses all of the information available, without making strong assumptions about the survey or its veracity.

## 2.1. Implementation

♣ **joey:** Describe Python implementation here ♣

## 3. Experiments

### 3.1. Simulating Faint Mira Variables

In this first experiment, we begin with a well-observed Mira variable from the ASAS Catalog of Variable Stars (ACVS, [Pojmanski et al. 2005](#)), ASAS 235627-4947.2. This star has a pulsation period of 266.6286 days and a Lomb-Scargle amplitude of 2.38 mag (V band). Note there are no non-detections in the ASAS light curve.

To simulate faint Mira stars from ASAS 235627-4947.2, we use the following procedure:

1. Convert the observed magnitudes,  $m_i$ , and errors,  $s_{m,i}$  to fluxes,  $f_i$  and flux errors,  $s_{f,i}$  (we assume a V-band zero-point of  $3.67 \times 10^{-9}$  erg/s/cm<sup>2</sup>/Å)
2. For a given flux dimming parameter,  $d$ , sample the new fluxes,  $\tilde{f}_i$ , from a Gaussian distribution centered around  $f_i/d$  with standard deviation sampled from the empirical ASAS distribution of  $s_{f,i}|f_i$ . ♣ **joey:** I found that relationship to be linear with a slope  $\approx 1$ , which is probably BS since the ASAS mag errors are all crap. So I added a reference value to each flux error to ensure a mag limit  $\approx 14.5$  mag. ♣
3. Fluxes that are not at least  $5\sigma$  above zero are denoted as non-detections and their flux estimates (and errors) are censored.

Folded light curves of ASAS 235627-4947.2, dimmed by four different values of  $d$ , are plotted in Figure 1.

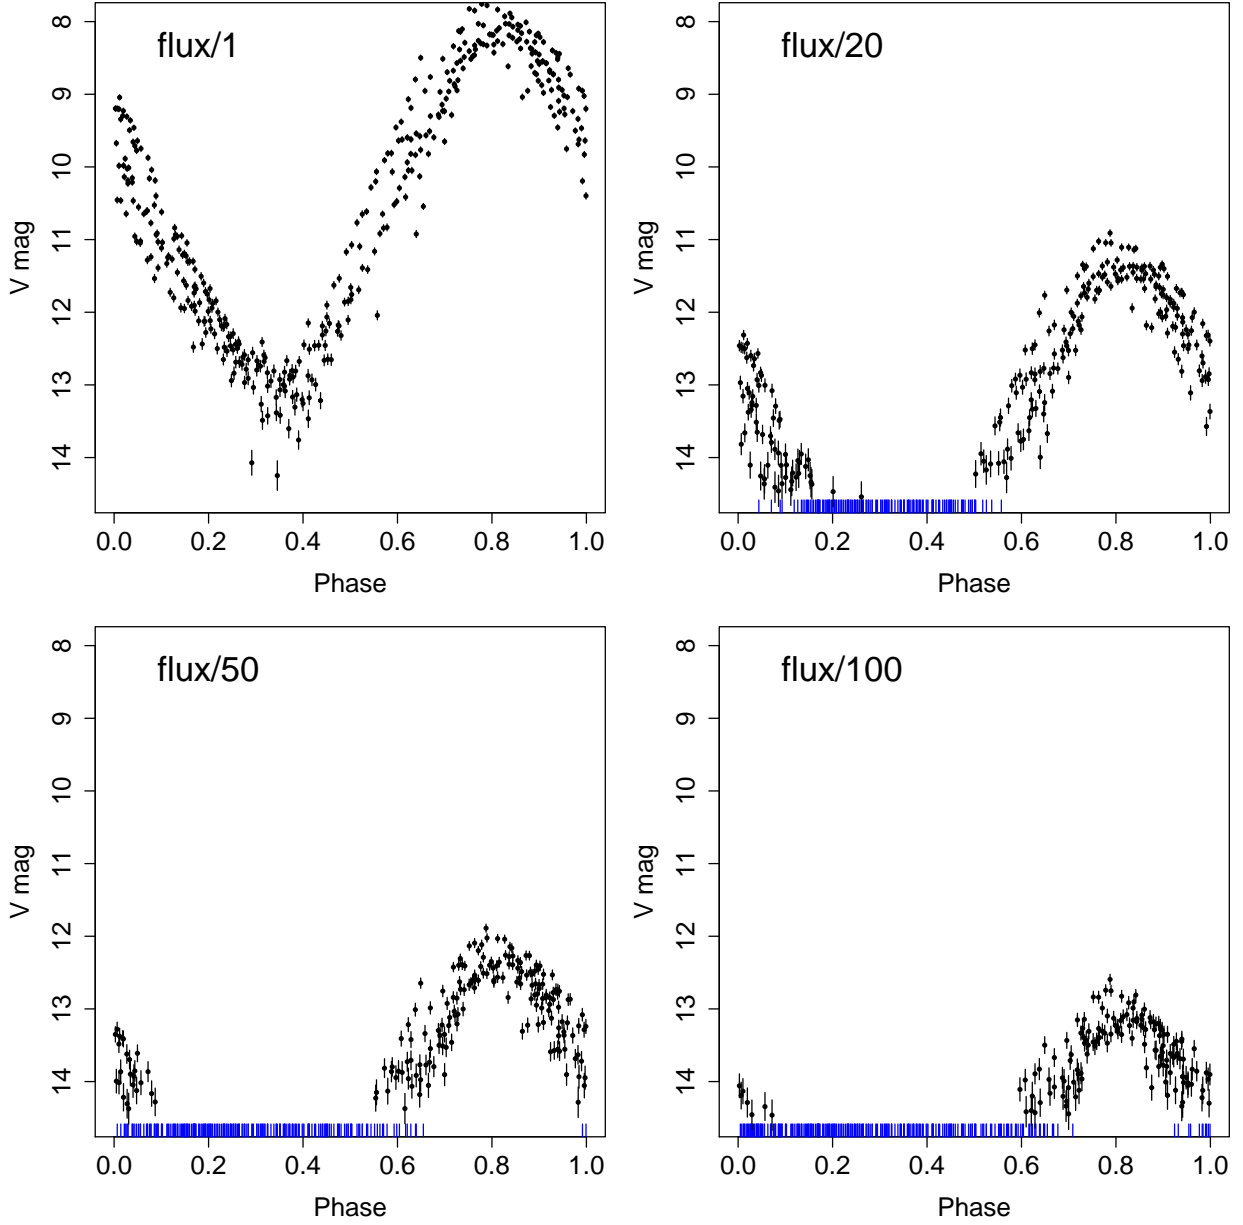


Fig. 1.— Folded light curves for data simulated from the Mira variable ASAS 235627-4947.2 using the prescription in §3.1. Blue tick marks along the bottom axis denote phases where there were non-detections. As the original flux measurements are dimmed by higher factors, the troughs of the sinusoidal light curve are censored, resulting in an incomplete view of the data.

### 3.2. Results of Fitting Miras

## 4. Results: ASAS Light Curves

We use the methodology of [Richards et al. \(2011\)](#) to select the top ASAS Mira and RR Lyrae, Fundamental Mode candidates. Using a posterior probability threshold of 0.8 gives us 1720 Mira and 1029 RR Lyrae candidates.

### 4.1. Mira Variables

Show P - A relationship before and after using the method

### 4.2. RR Lyrae Variables

Show P - A relationship before and after using the method

There is an RRL P - A relationship (strong linear anti-correlation)

## 5. Discussion

What people can do, starting from the data-taking procedure.

limitations

## REFERENCES

- Pojmanski, G., Pilecki, B., & Szczygiel, D. 2005, *Acta Astronomica*, 55, 275
- Richards, J. W., et al. 2011, *ApJ*, 743